

Normalization: → Organizing data into multiple related tables
(also known as Schema Refinement)
→ Minimize Data Redundancy (Not elimination)
→ eliminates Anomalies like Insertion, Deletion, Updation.

1NF

- each cell must have single value
- column should have values of same datatype
- Each column should have unique name.
- Order of data doesn't matter.

2NF

- must be in 1NF
- NO partial dependency
 - A column depending only on one field
- can be eliminated by moving the field into another table.

3NF

- must be in 2NF
- NO Transitive dependency
 - $A \rightarrow B \rightarrow C$
Transitive
- Create New table.

Decompositions: Decomposition is done when table is not in appropriate normal form.

- Breaking into multiple tables.
- If no proper decomposition is done, it leads to loss of info.
- used to eliminate problems like bad design - anomalies, inconsistencies & redundancy.

Types:

1. Lossless decomposition:

- After decomposition, information is not lost.
- guarantees that join of relation results in same relation before decomposition.

2. Lossy: If joined after decomposition, then data may be lost.

Properties:

1. Dependency preserving:

- functional dependencies b/w attributes in a relation are preserved when transforming the table from one NF to another.

2. Decomposition must be lossless.

3. Lack of Data redundancy.

Candidate key: It is the minimal set of attributes whose attribute closure is set of all attributes of relation is called the candidate keys of relation.

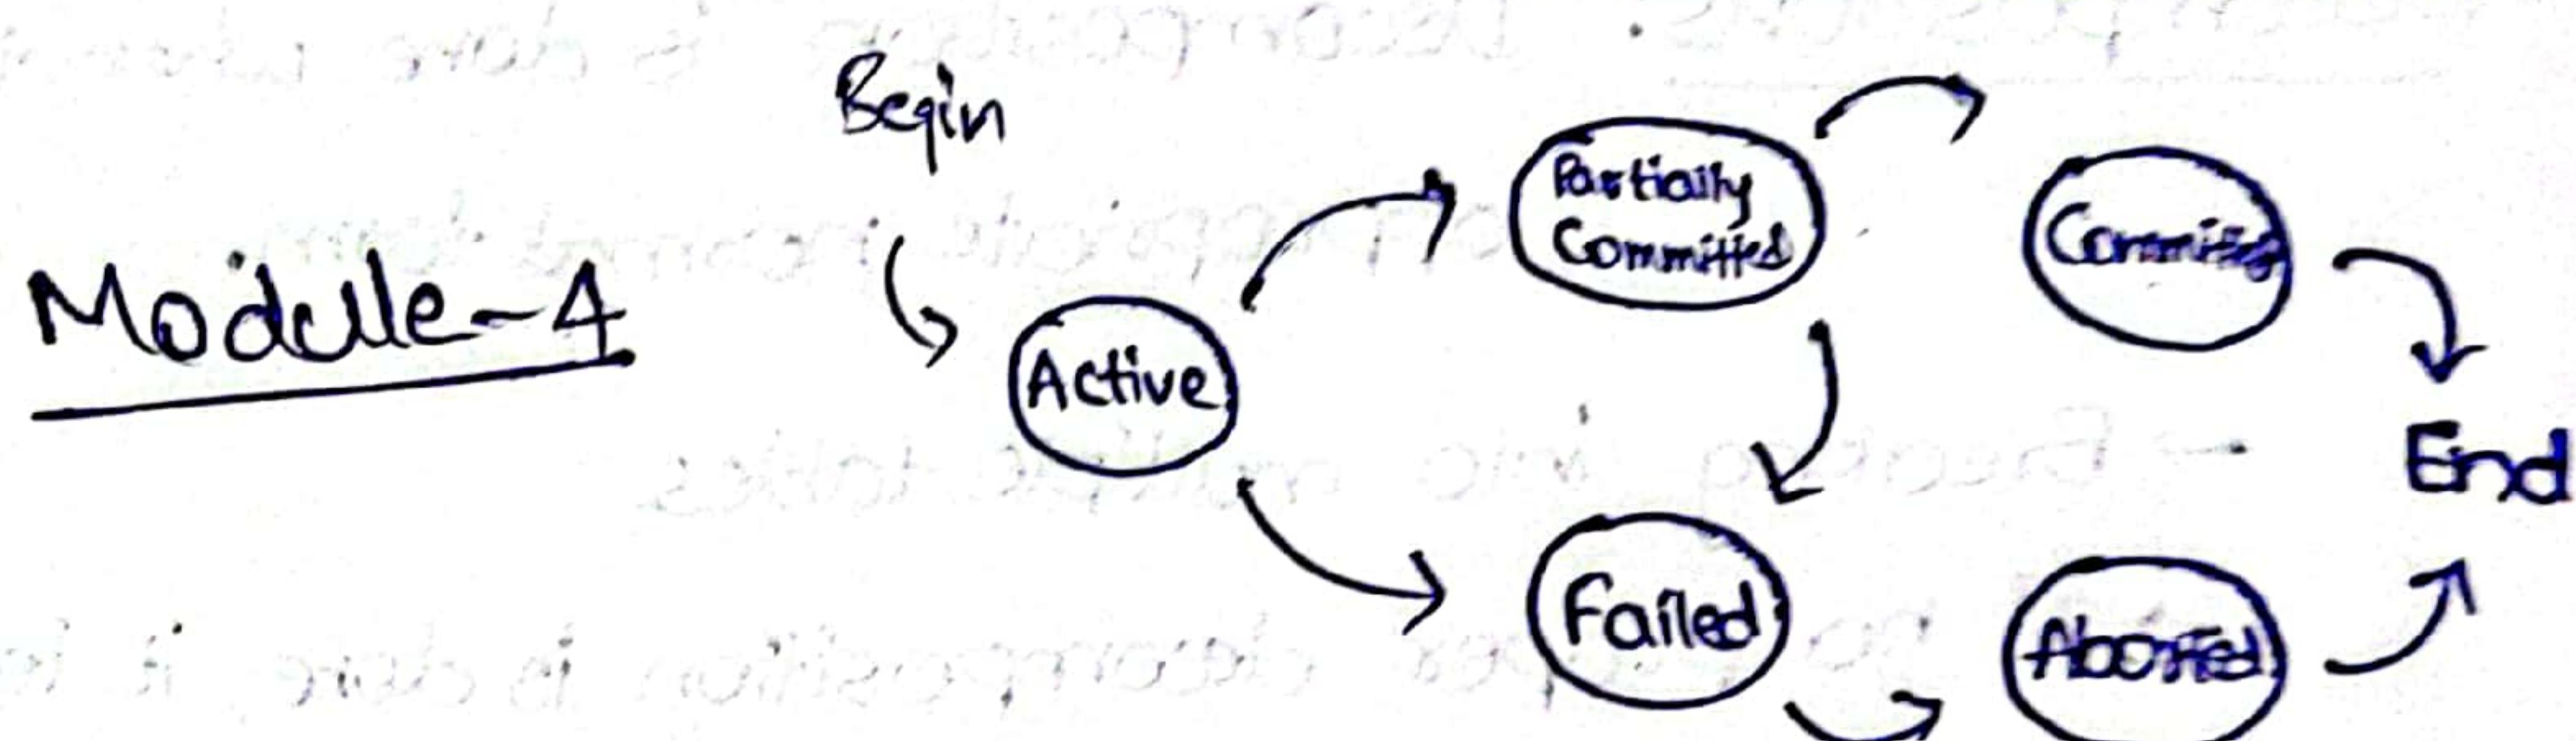
Primary key

- Identify each record in DB
- NO NULL Values
- Only one P.K.
- Indexed automatically
- Cannot be deleted
- Create parent-child relationship

foreign key

- links two tables together.
refers to Primary key of 2nd table.
- Can have NULL values
- > 1 can be acceptable.
- Indexed Manually (can)
- Can be deleted
- Create parent-child relationship

Phases of Transaction



ACID Properties:

1. Atomicity :- (also "all or nothing rule")

Either the transaction takes place or doesn't happen at all. There's no midway like partial transaction.

Two operations : Abort & Commit.

2. Consistency:- DB must be consistent before & after the transaction.

Refers to correctness of DB.

3. Isolation:- Multiple transactions occur independently with no interference.

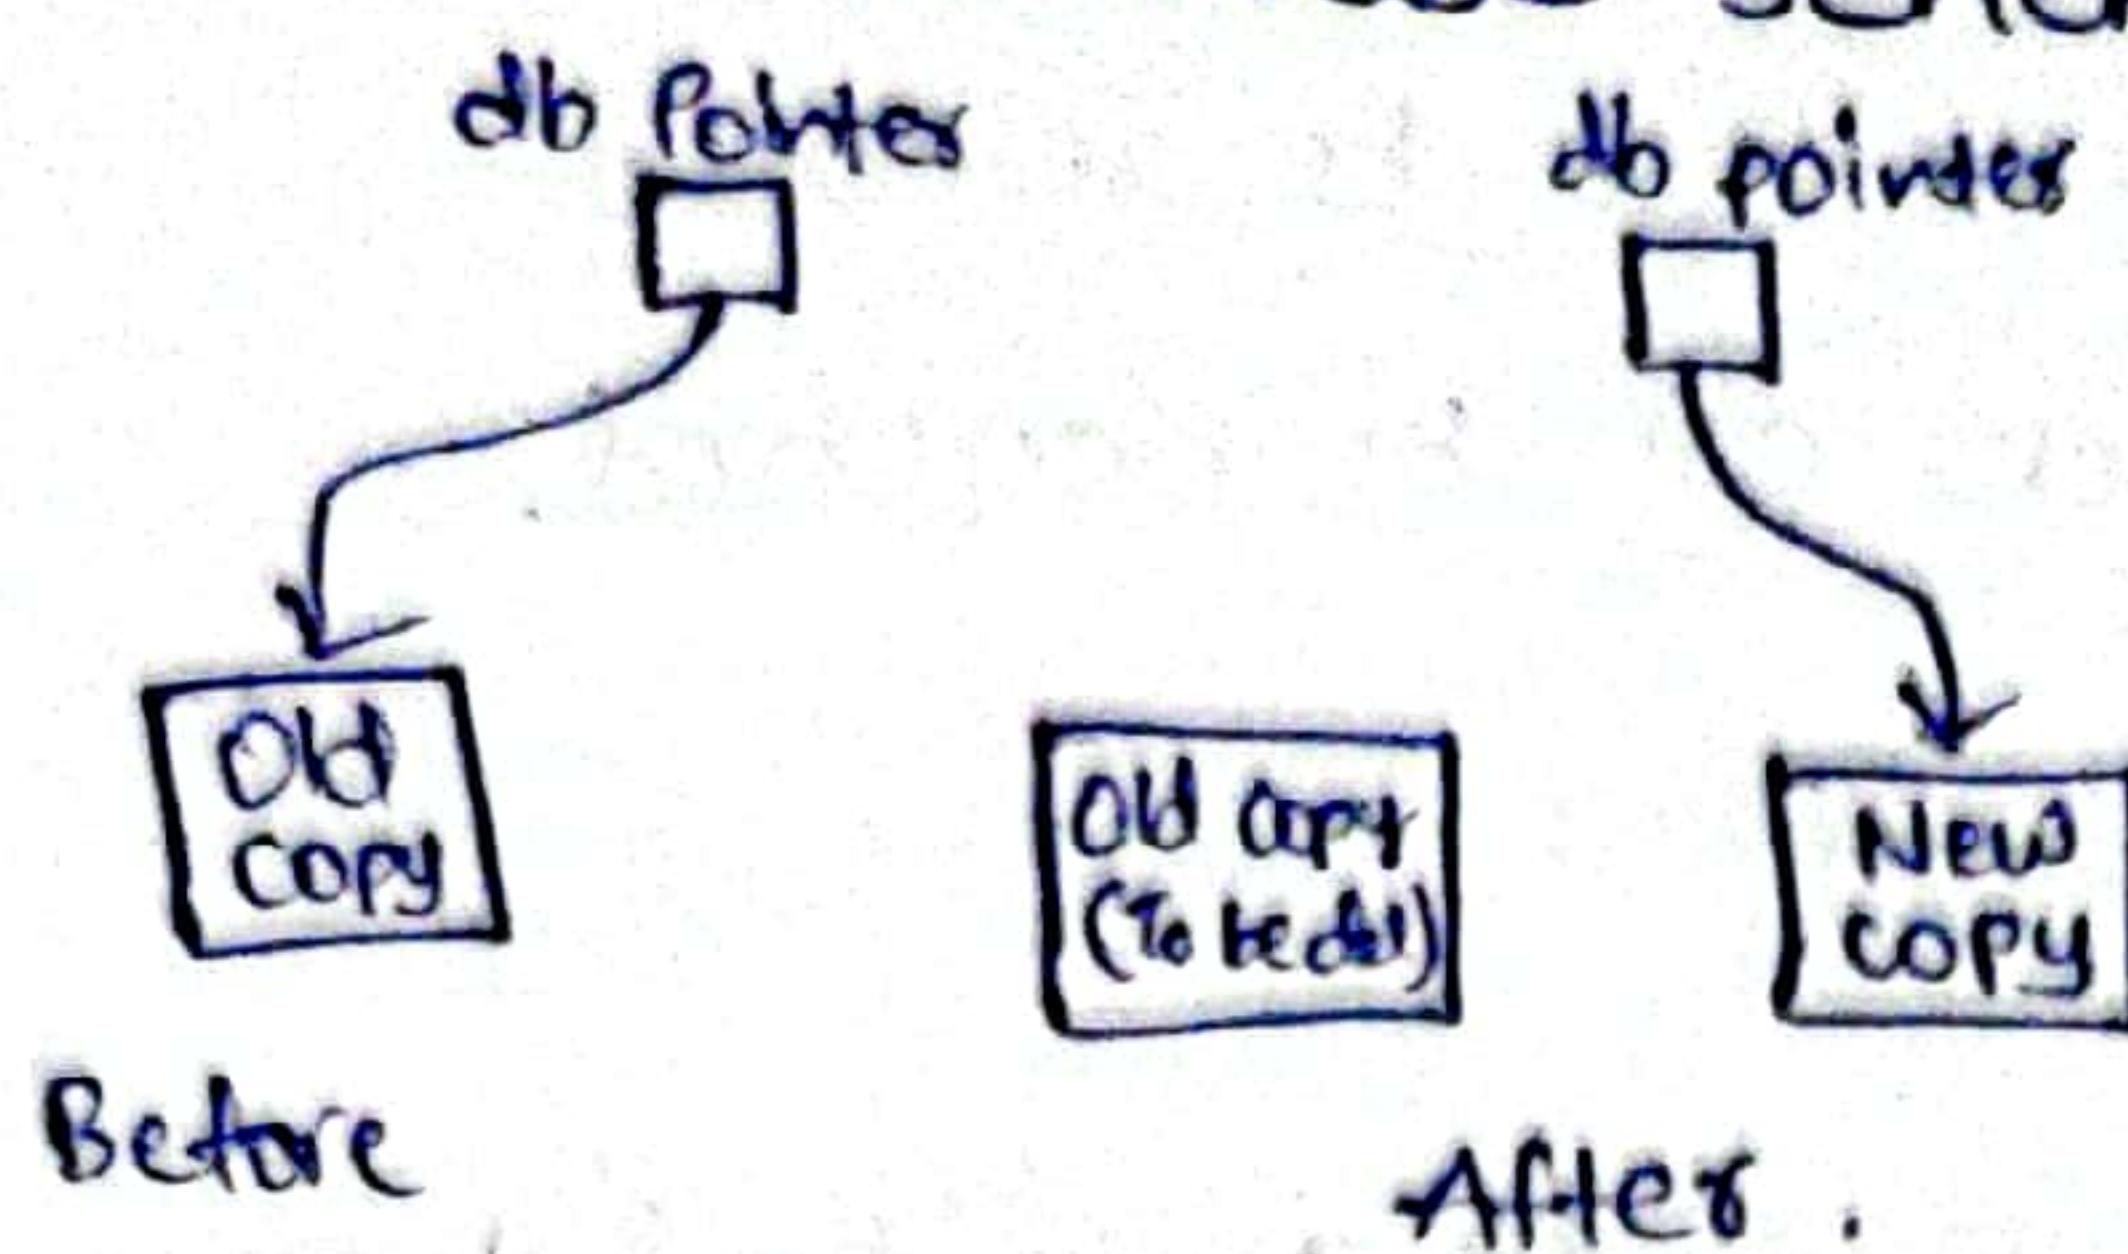
Changes occurring in particular transaction will not be visible to any other transaction until it occurred.

4. Durability:- After transaction, updates & modifications to DB are stored in & written to disk & they persist even if system failure occurs. Thus, effect of transactions are never lost.

Implementation of Atomicity & Durability

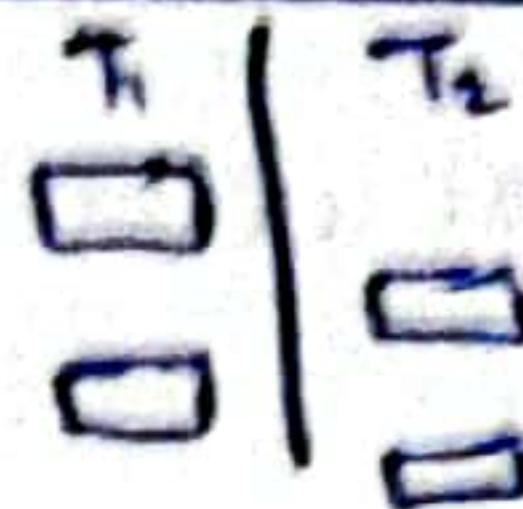
→ Recovery-management component.

→ Shadow-database scheme



Interleaved manner
↓
Alternate fashion.

Concurrent execution of Transaction

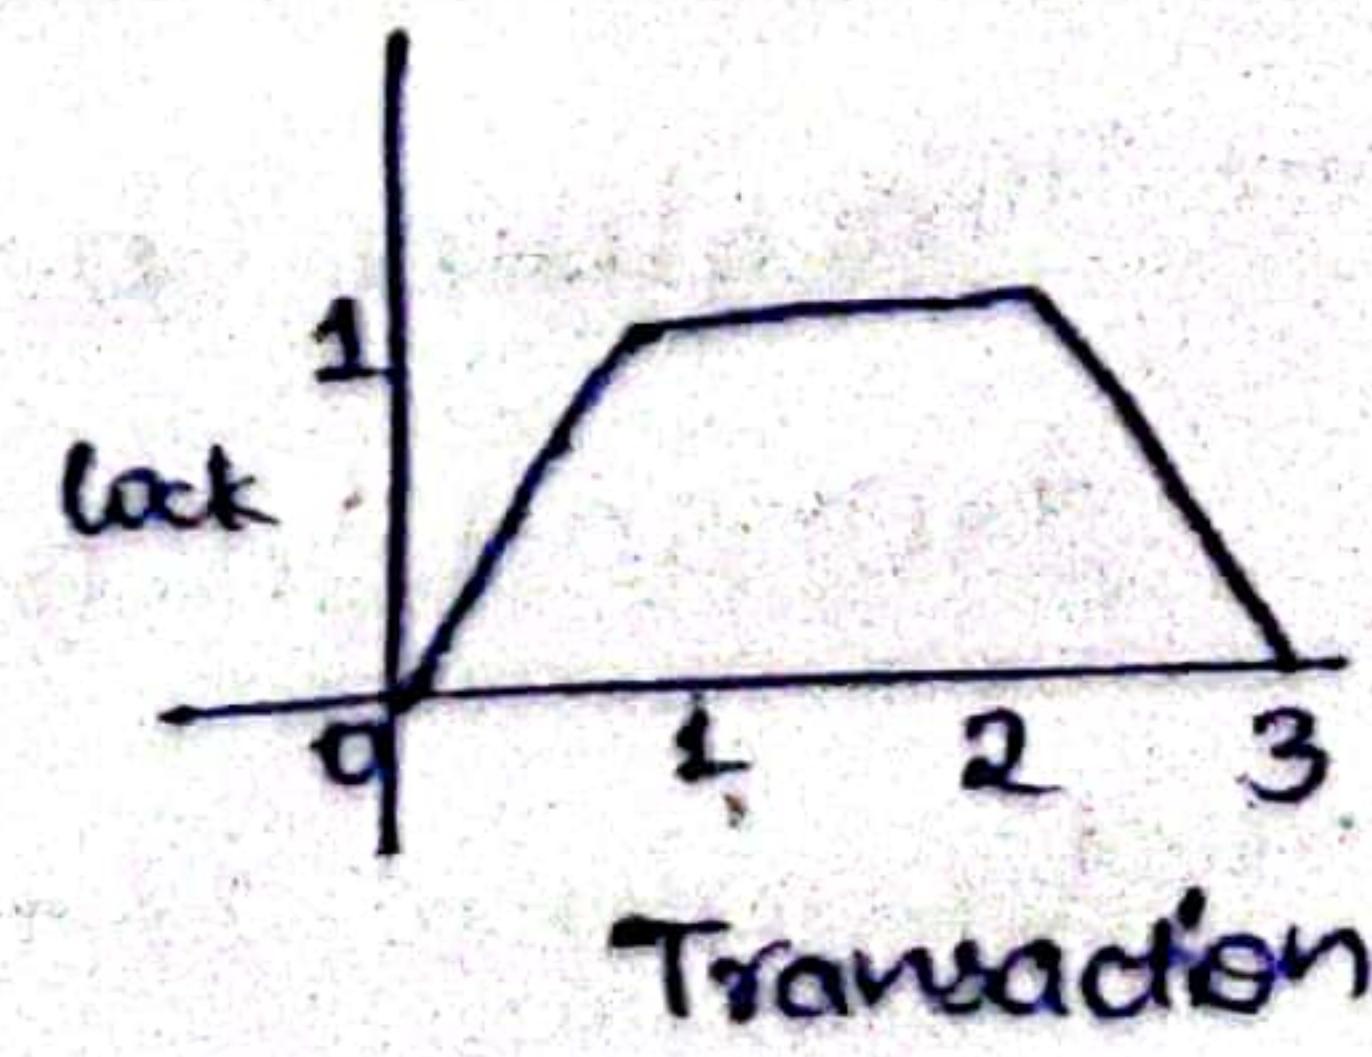
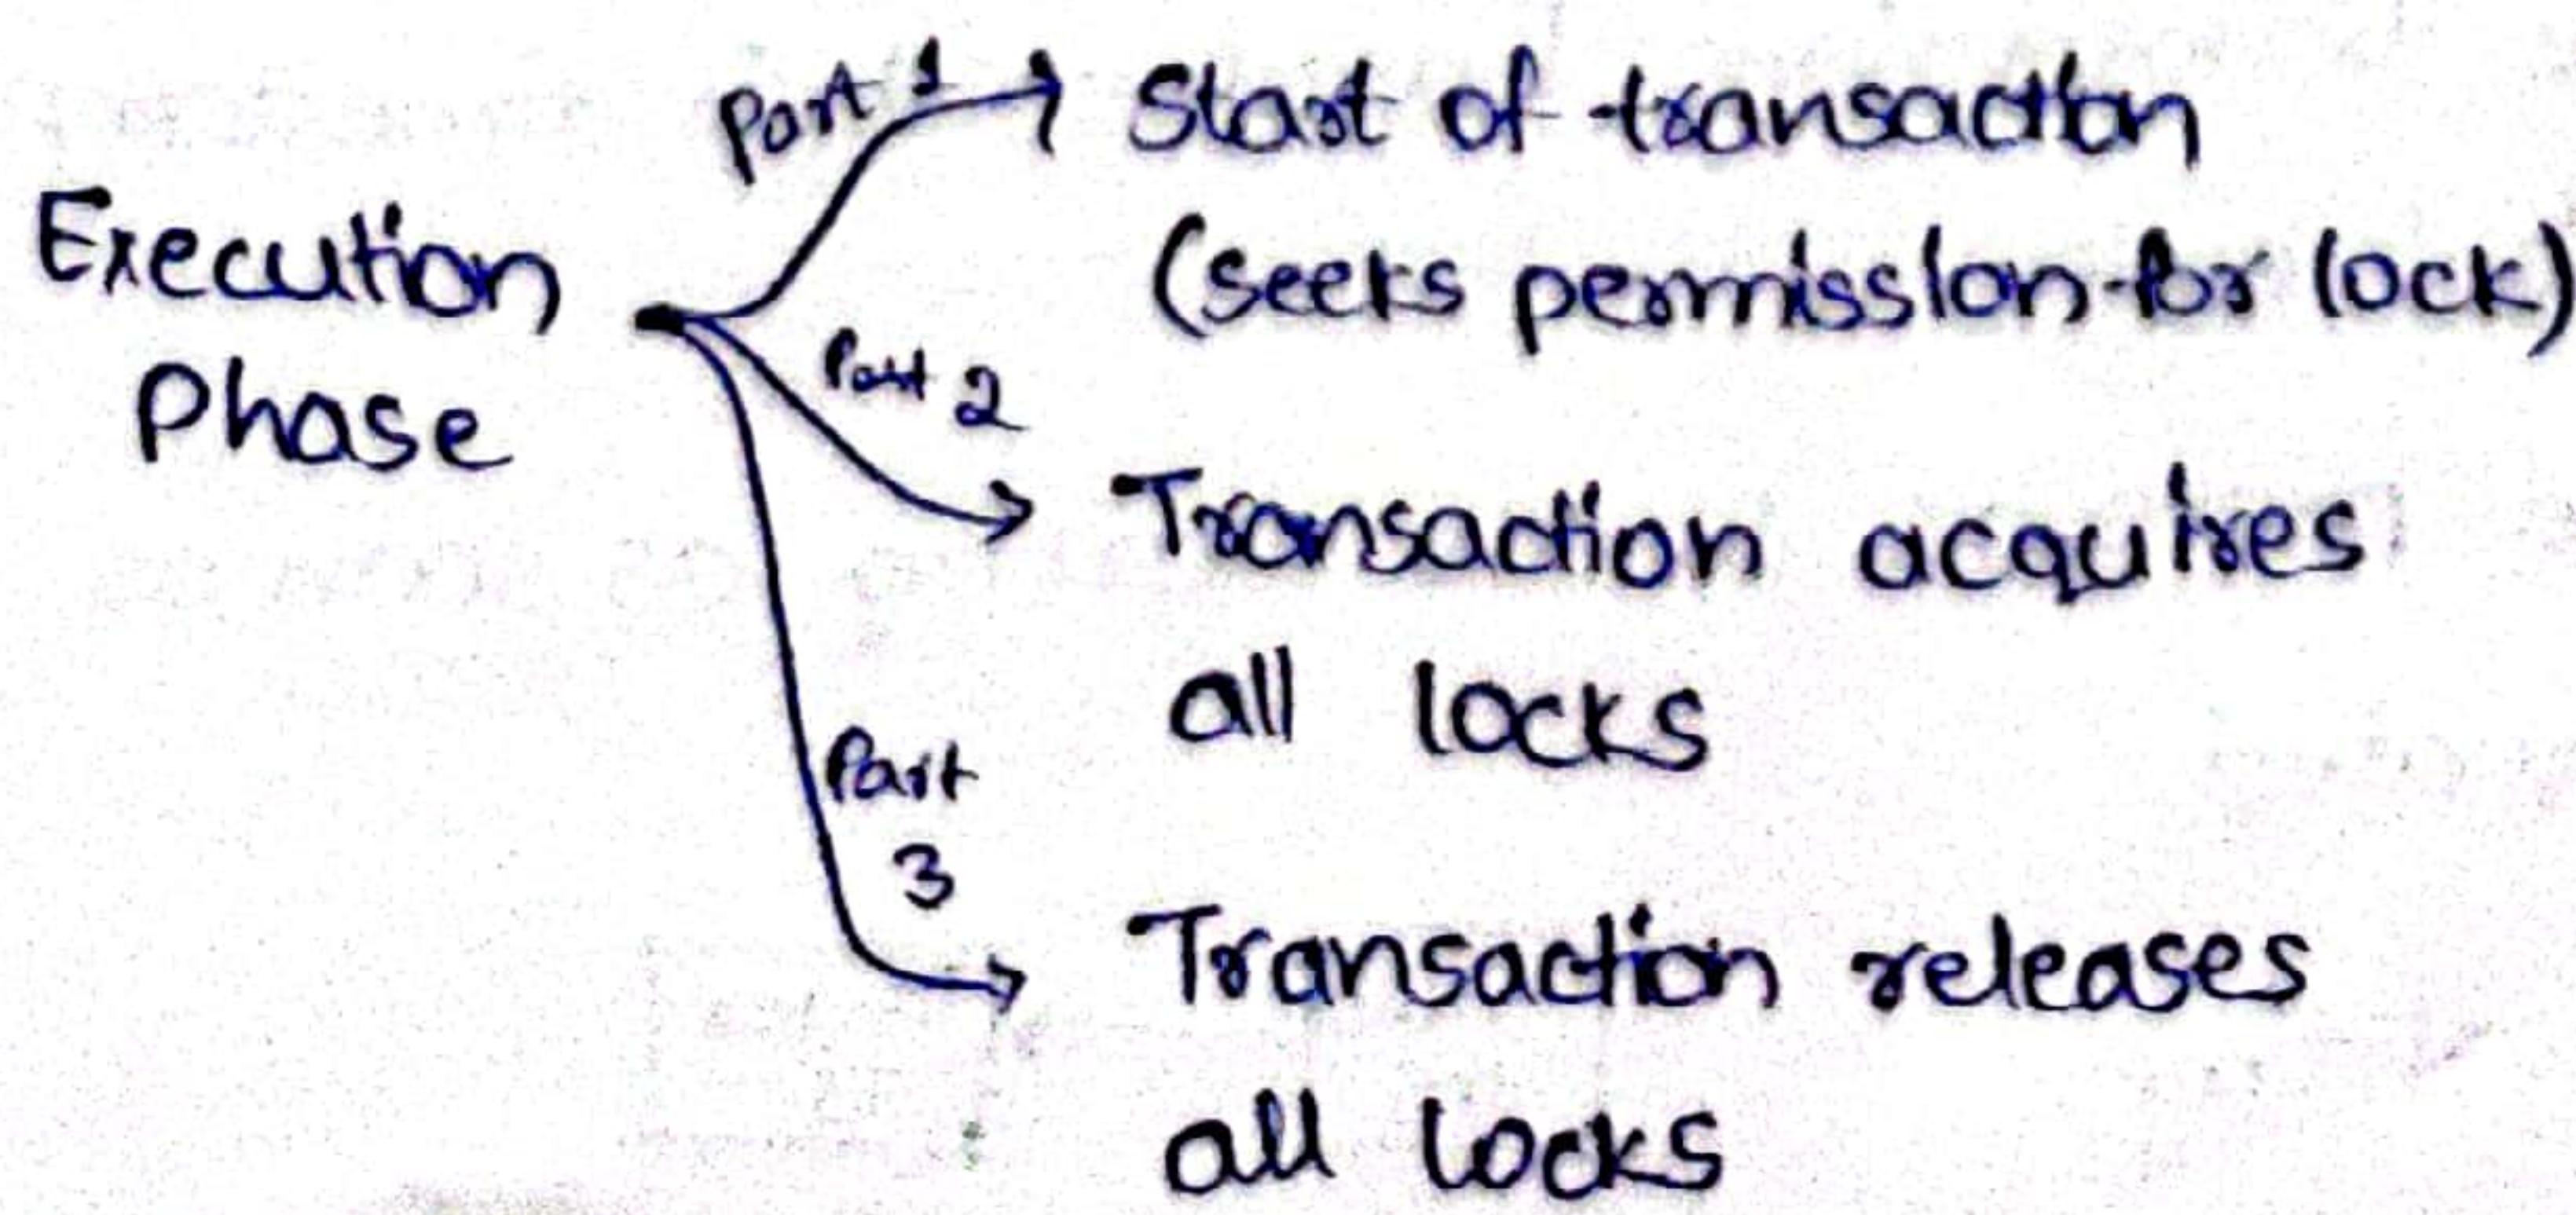


- Multi user system - multiple users - same database - sometime
- Simultaneous execution in interleaved manner.
- No operation should affect other executing operations - maintaining the consistency.
- DB Transactions has 2 main operations : R/W. If ~~not~~ performed in interleaved manner , the data may become inconsistent.

Serializability:

- Serial schedule - executed in series (one after other)
- Serializability is used to find correct non-serial schedules in DB.
- Maintains consistency & relates to isolation features of transaction.
- Serializability decides if an interleaved non-serial schedule is serializable or not.

Two-Phase Locking Protocol (2PL)

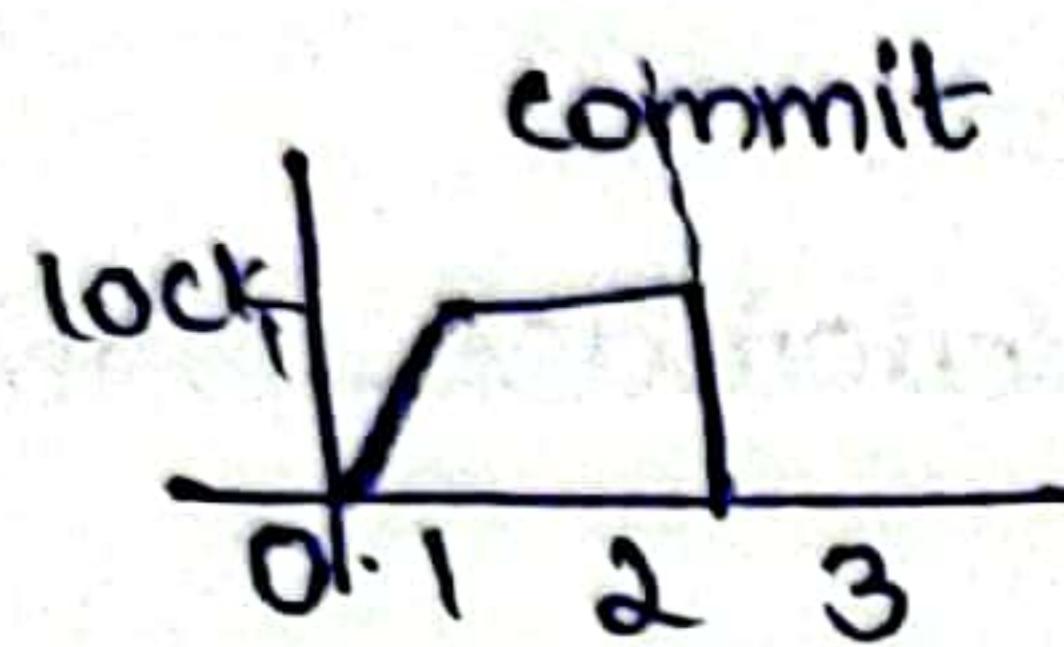


Phases of 2PL:

1. Growing Phase: New lock on data can be acquired
no lock can be released.
2. Shrinking phase: existing locks can be released
no new locks can be acquired.

Strict 2PL:

- Same as 2PL but holds all locks till the end of transaction.
- Locks will be released instantly after committing transaction.



Concurrency control protocols:

- To enforce isolation
- To preserve DB consistency
- To resolve conflicts of RW, WR, WW.

1. Lock Based Protocols:

- A transaction requires appropriate lock to read & write data.

Shared lock: Read-Only lock.

↳ Can be shared b/w the transaction as it cannot be written.

Exclusive lock: R/W lock.

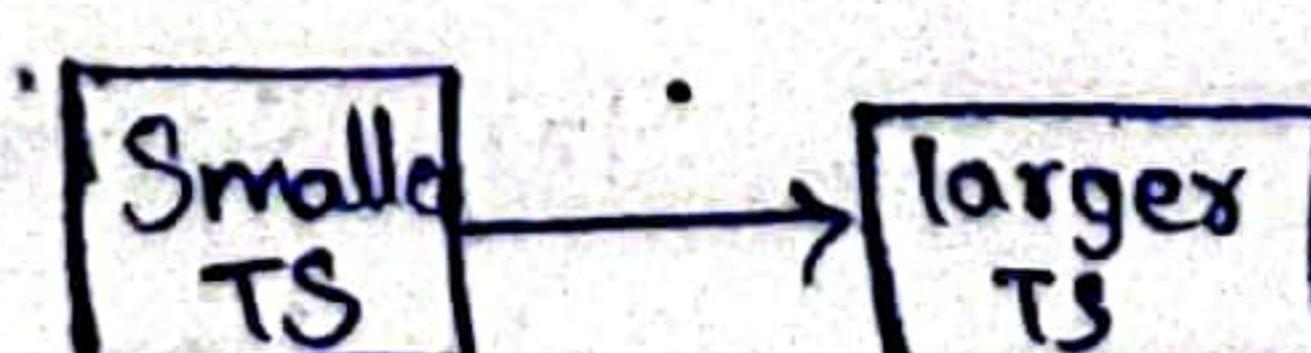
↳ Allocated only to single transaction at particular instance of time.

2. Timestamp based protocols:

- used to order the transactions based on their timestamps.
(in the chronological order)

→ Timestamp based protocols starts working as soon as transaction is created.

- Priority $\propto 1/\text{arrival time}$.



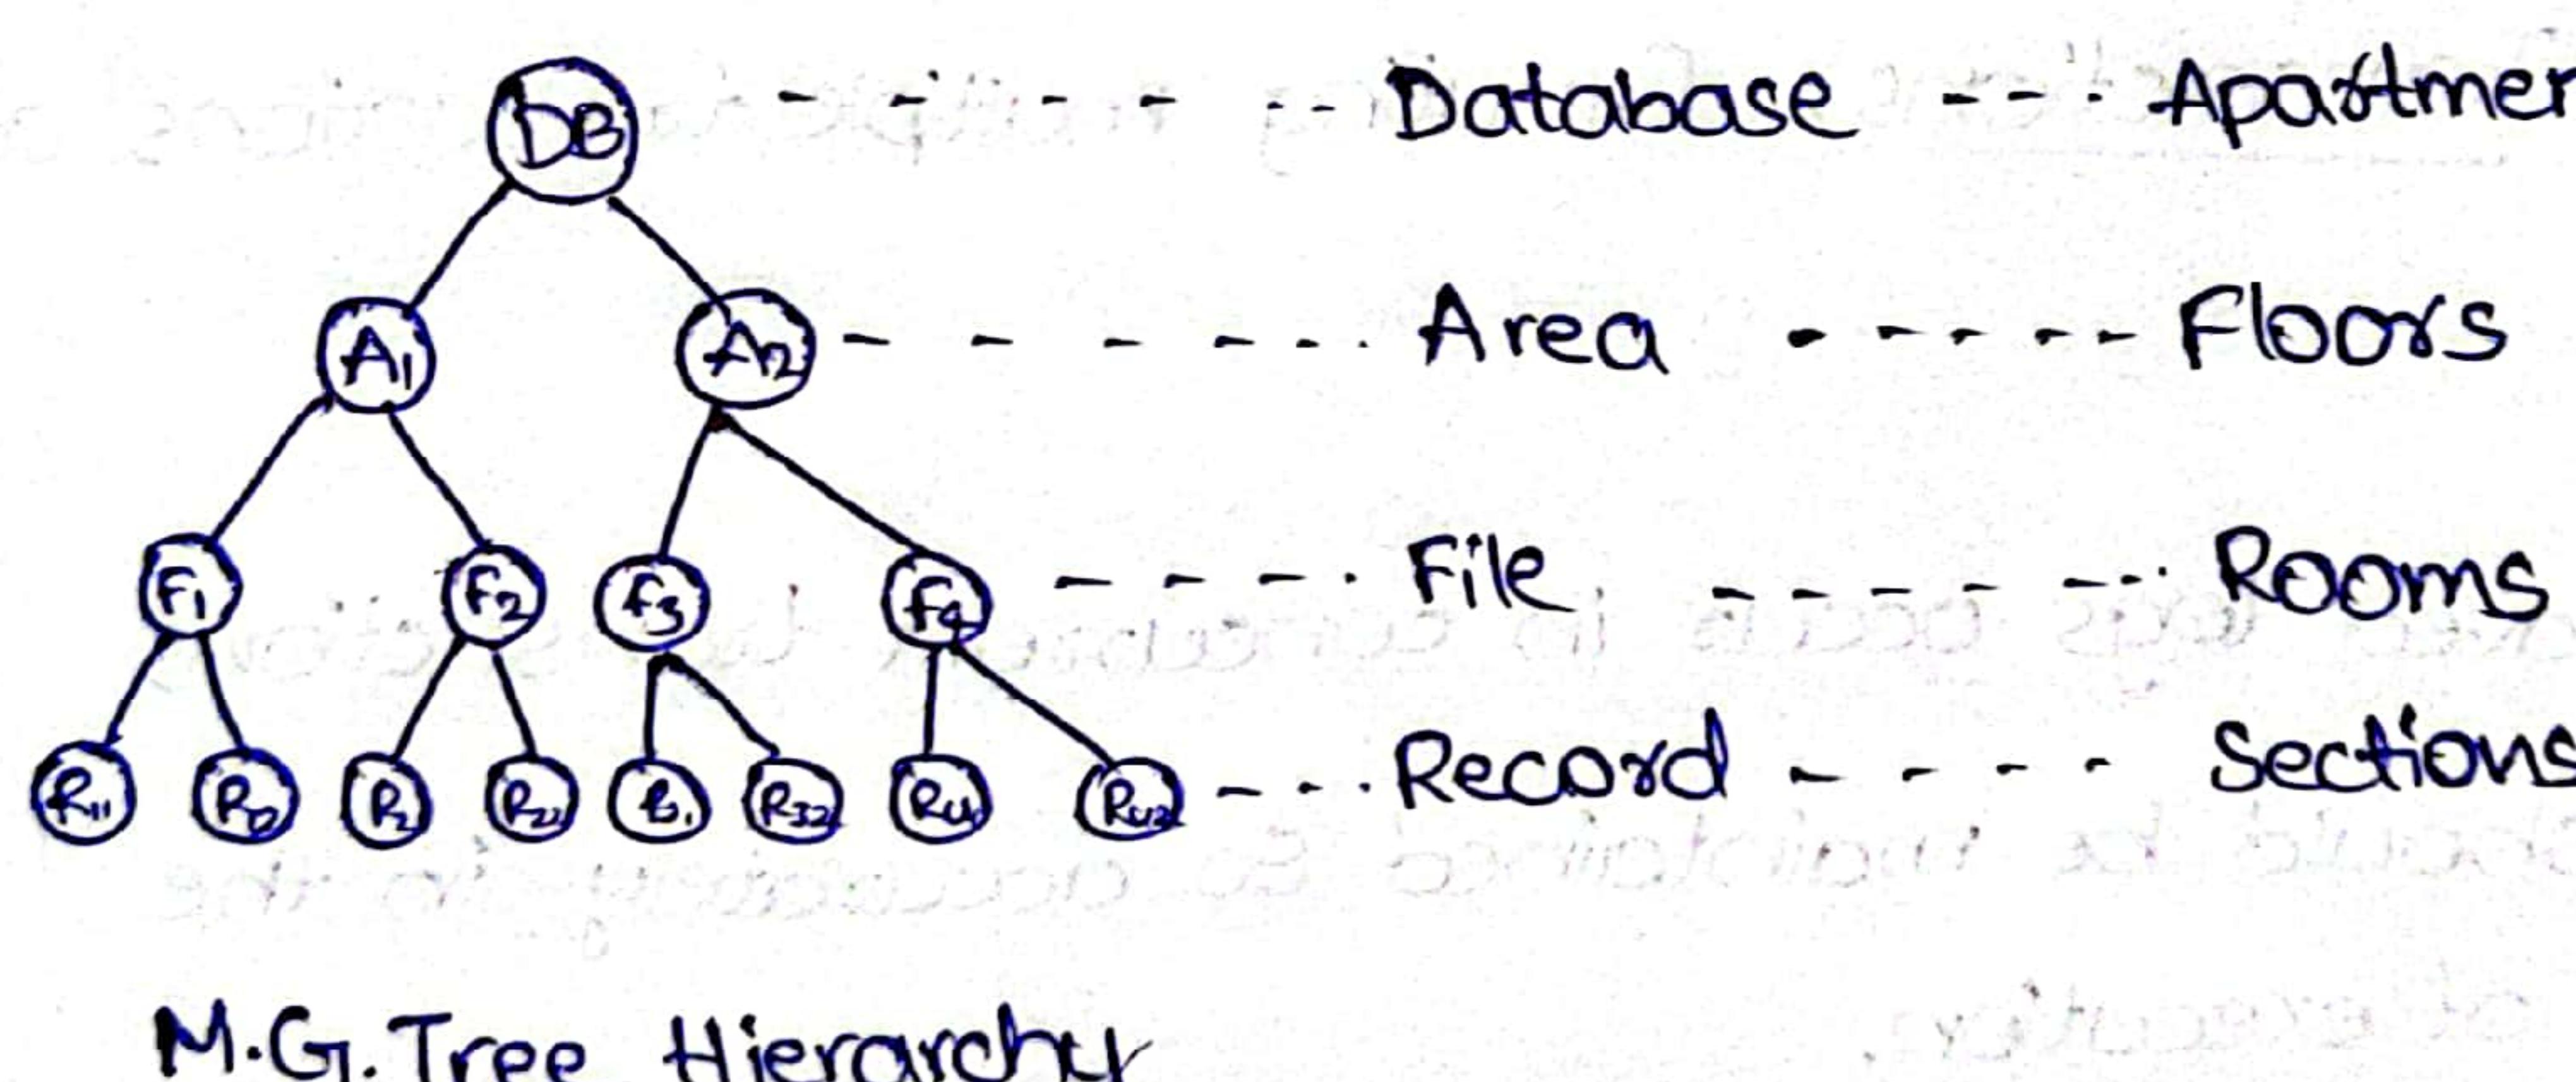
3. Validation based protocols: (also Optimistic concurrency control)
↓
Very less interference occurs
technique

3 Phases of Execution

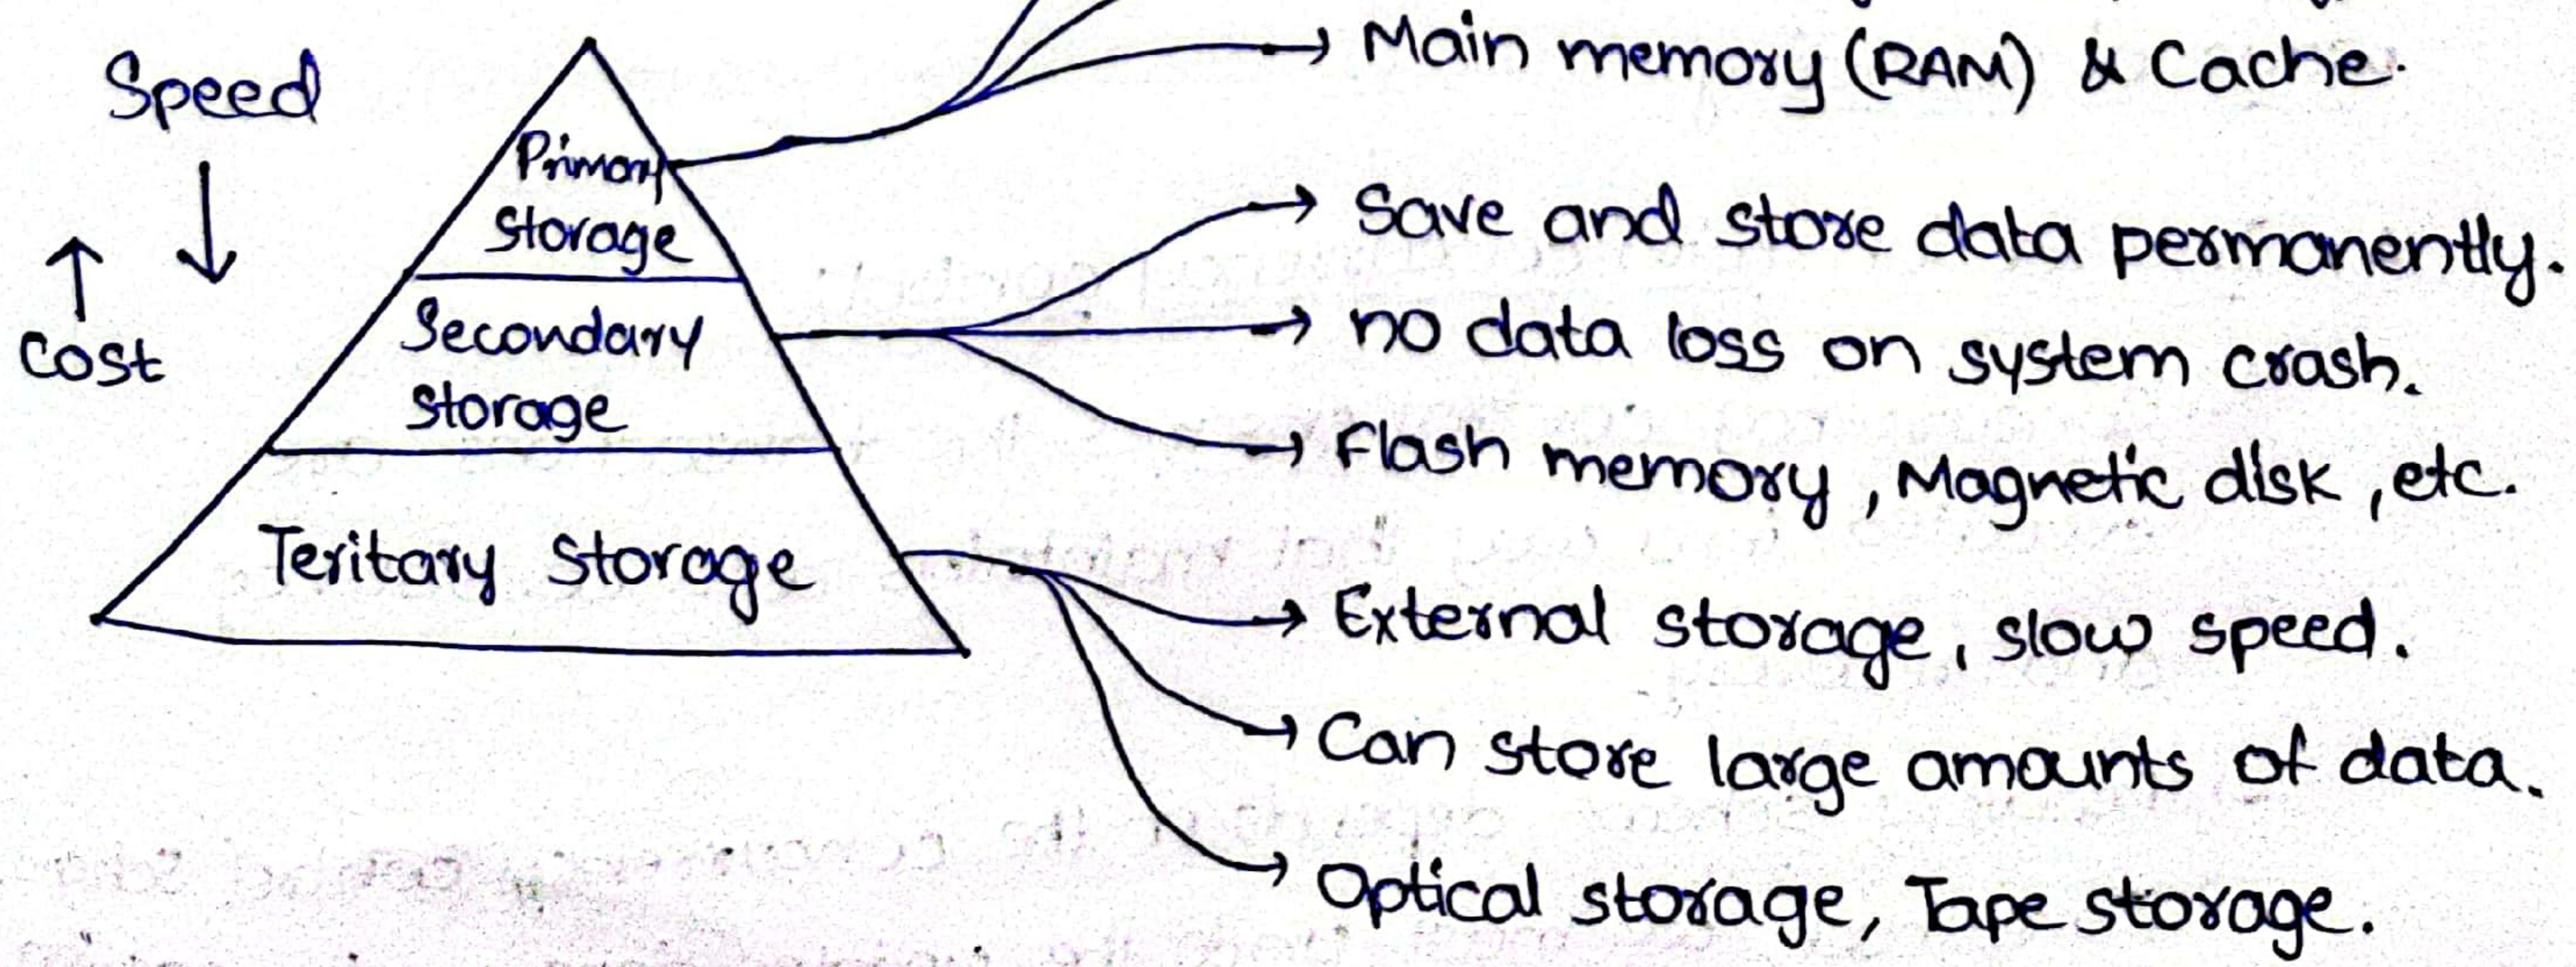
- 1 → Read Phase : → reads various data items and stores in temporary variables.
→ Can modify temp. variables without changing actual DB.
- 2 → Validation phase : → Temp. variables will be validated with actual DB to check with serializability.
- 3 → Write phase : → After validation, temp results are written to DB or transaction will be rolled back.

Multiple Granularity: (lol... it's a protocol) [enhances concurrency]

- Hierarchically breaking up the DB into blocks which can be locked.
- Maintains the track of what and how to lock.
- locking and unlocking of a data item made easy.
- Represented as a tree.



Storage structure:

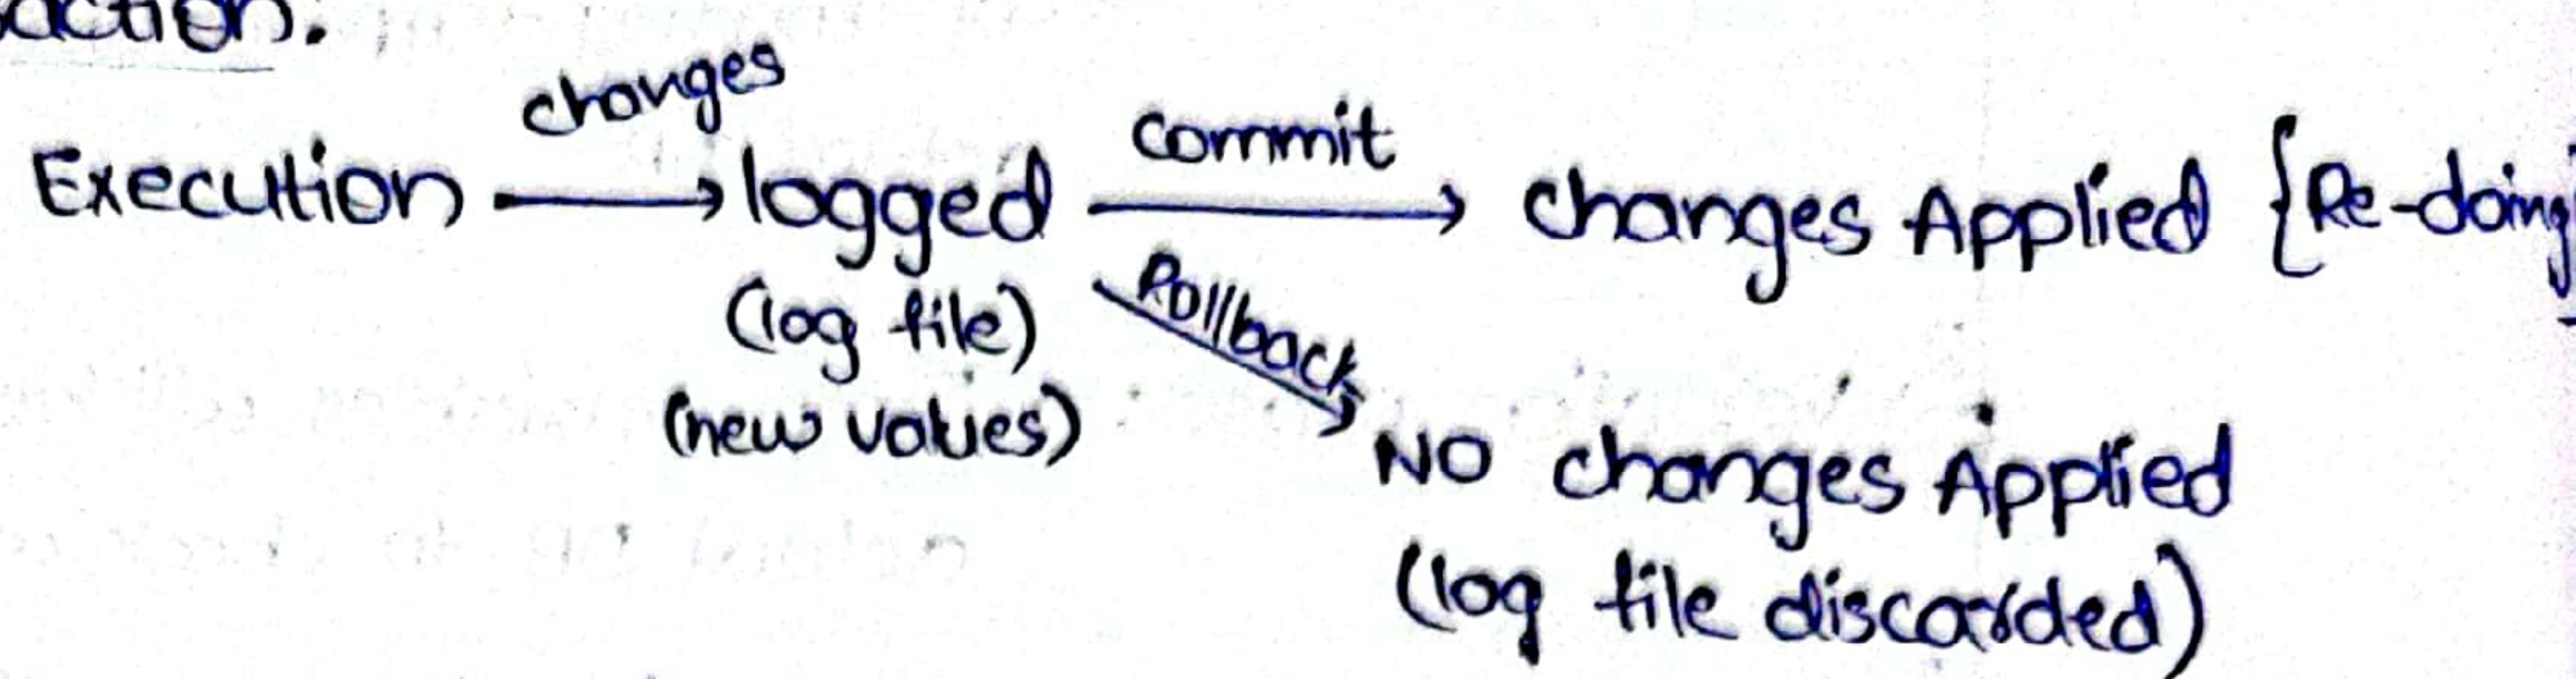


Database Modification:

1. Deferred update: (also NO-UNDO/REDO technique)

- Maintains log files of transactions.
- used to recover transaction failures occurred unexpectedly.

→ Transaction:



2. Immediate update (also UNDO/REDO technique)

- Same as previous but log files contain both old as well as new values.
- If rollback, old values of log file are stored in the DB
- used for recovery of transaction failures that occur due to power, memory, OS failures.

Concurrent Transactions: Executing multiple transactions at a time.

Recovery:

- Interleaved logs occur in concurrent transactions
- logs should be maintained so accurately in the order of execution.
- During recovery, it'd be very difficult for the recovery system to backtrack and recovering the logs.

ways of recovering:

1. Interaction with concurrency control:

- Concurrency control ensures the transactions are executed in a way that maintains data consistency and recovery.

- Recovery scheme depends on the concurrency control scheme. To rollback, we must undo the updates done by transaction.

2. Transaction Rollback:

- use logs to rollback a failed transaction.
- System scans the log backwards a failed transaction, for every log record found in the log, the system restores the data item.

3. Checkpoints:

- Saving a snapshot of the application state.
- Checkpoints are the predefined points in time where the DBMS records the state of DB & transaction logs.
- Creation: DBMS creates a checkpoint by writing all the modified data from memory to disk, and recording the checkpoint in the logs.
- Recovery: System starts applying changes from logs starting from most recent checkpoint rather than from beginning, reducing recovery time.

4. Restart recovery:

- This process aims to bring the system back to its normal operation after a failure. It involves replaying the transactions that were committed but not fully yet before the failure.
- Redo, Transaction commit, Resuming normal operations.

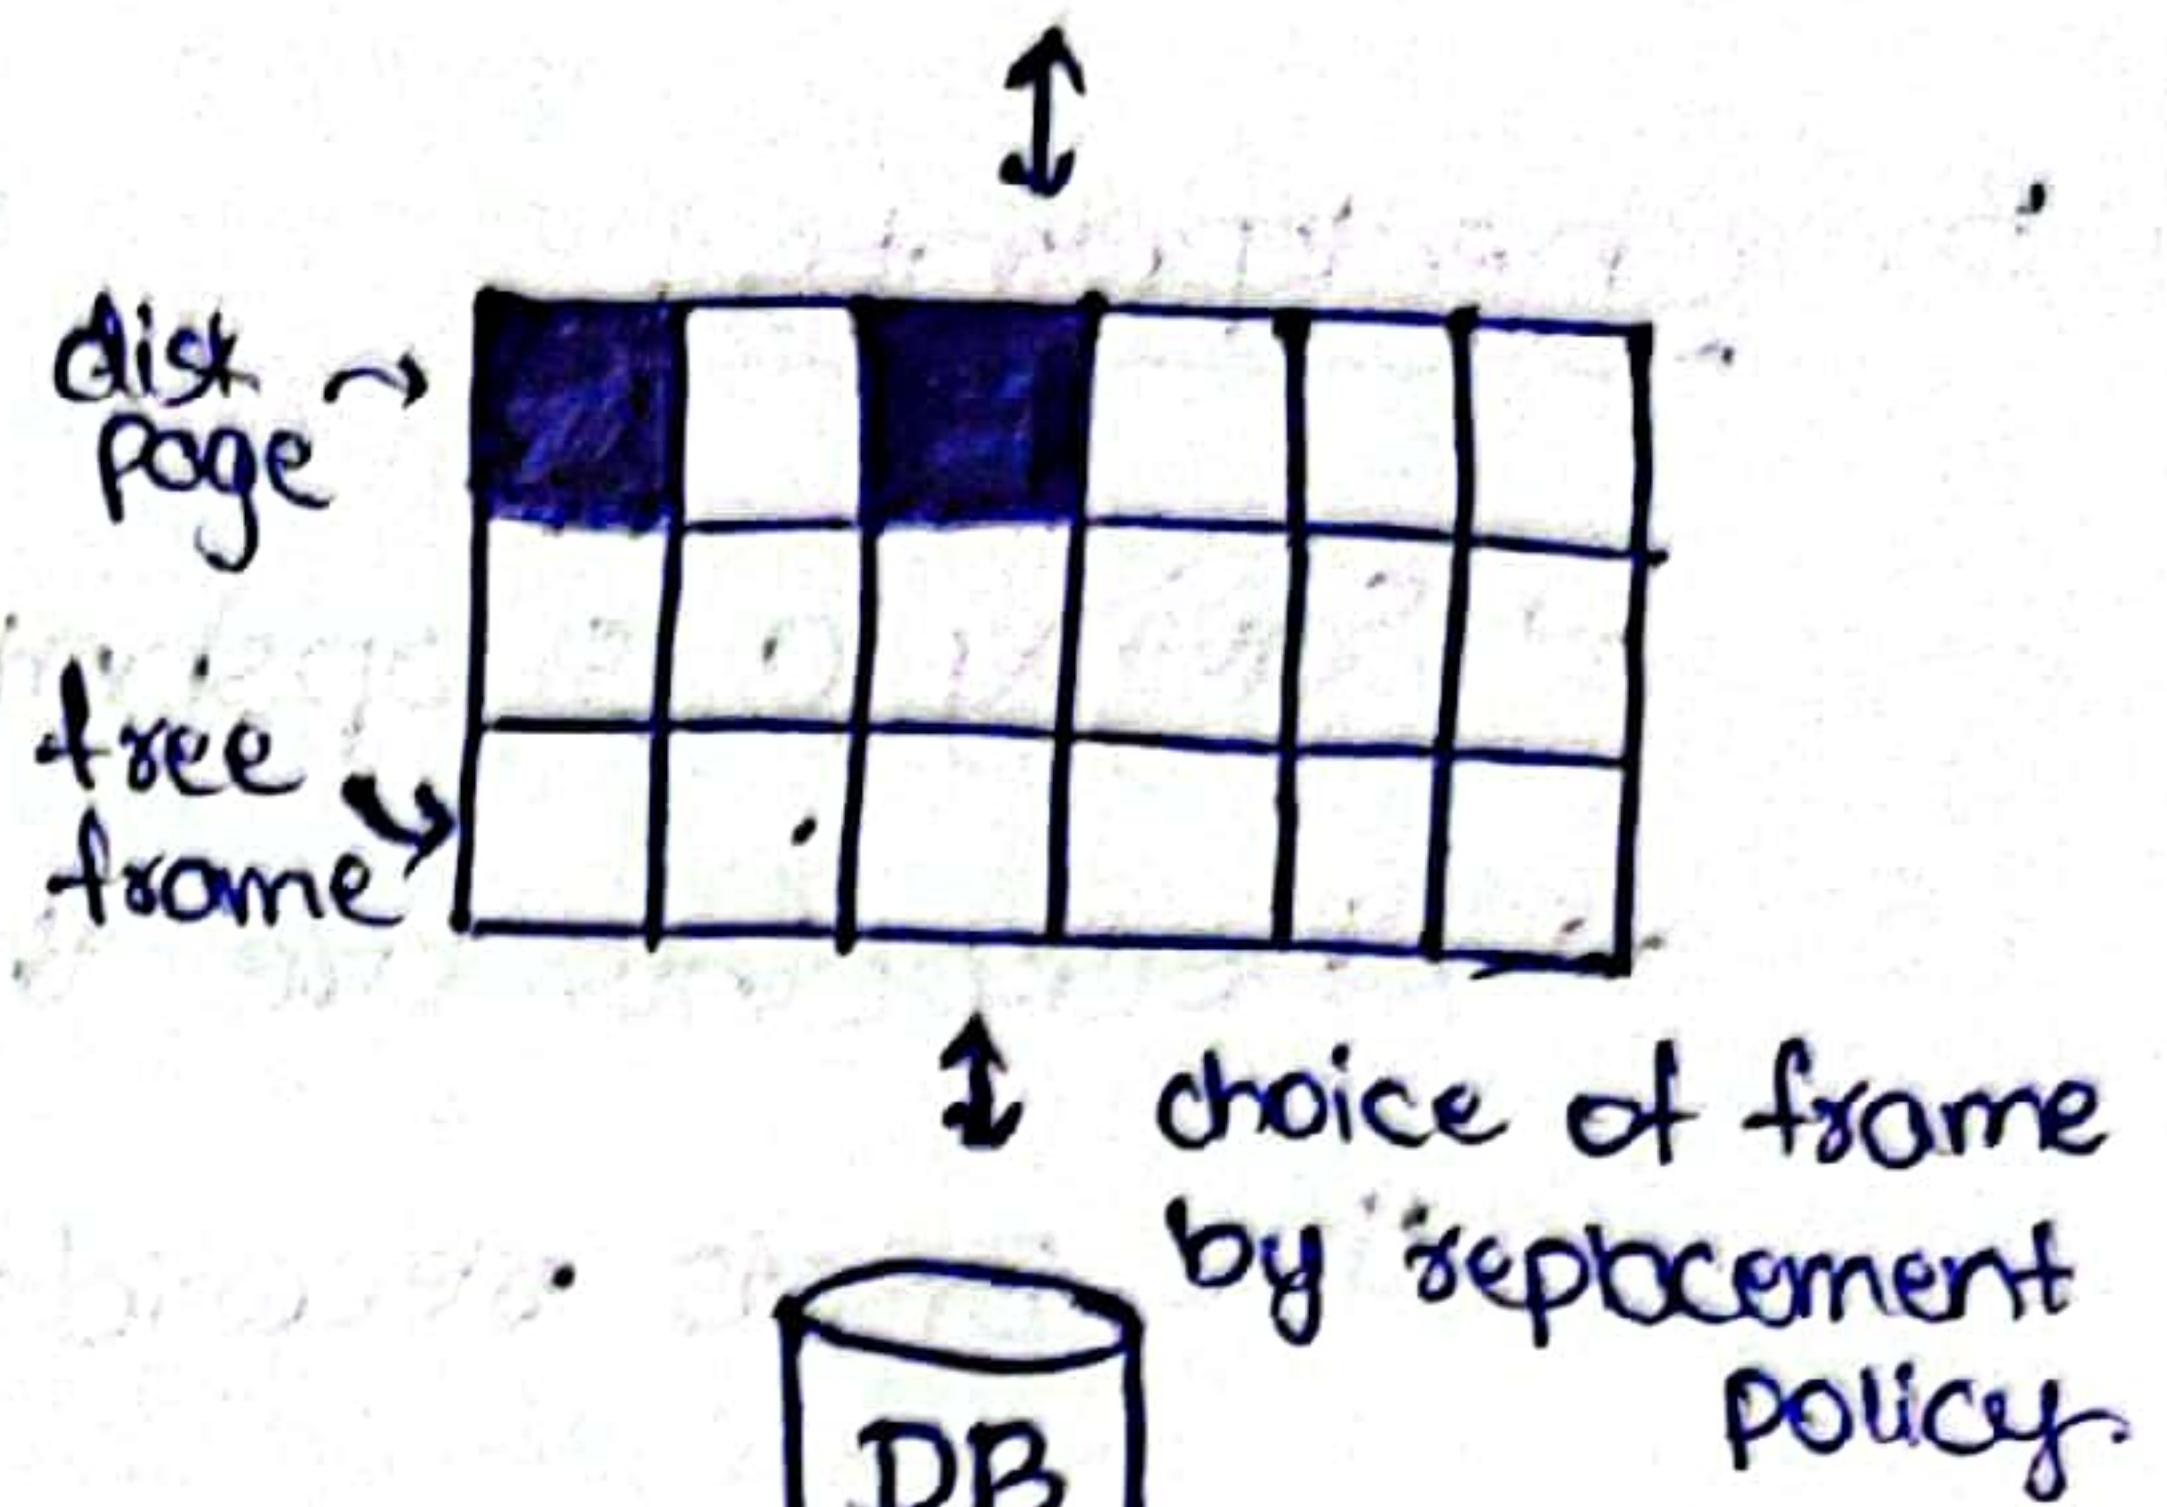
Buffer Management:

- Buffer Managers - allocates space to the buffer to store data into buffer.
- User requesting → If block available → Buffer Manager provides the block address in Main memory.
 - a particular block → if block not available → Buffer Manager allocates the block in the buffer.

→ No free space → existing blocks were thrown from the buffer for space + these were written back into disk.

[location in Main memory]

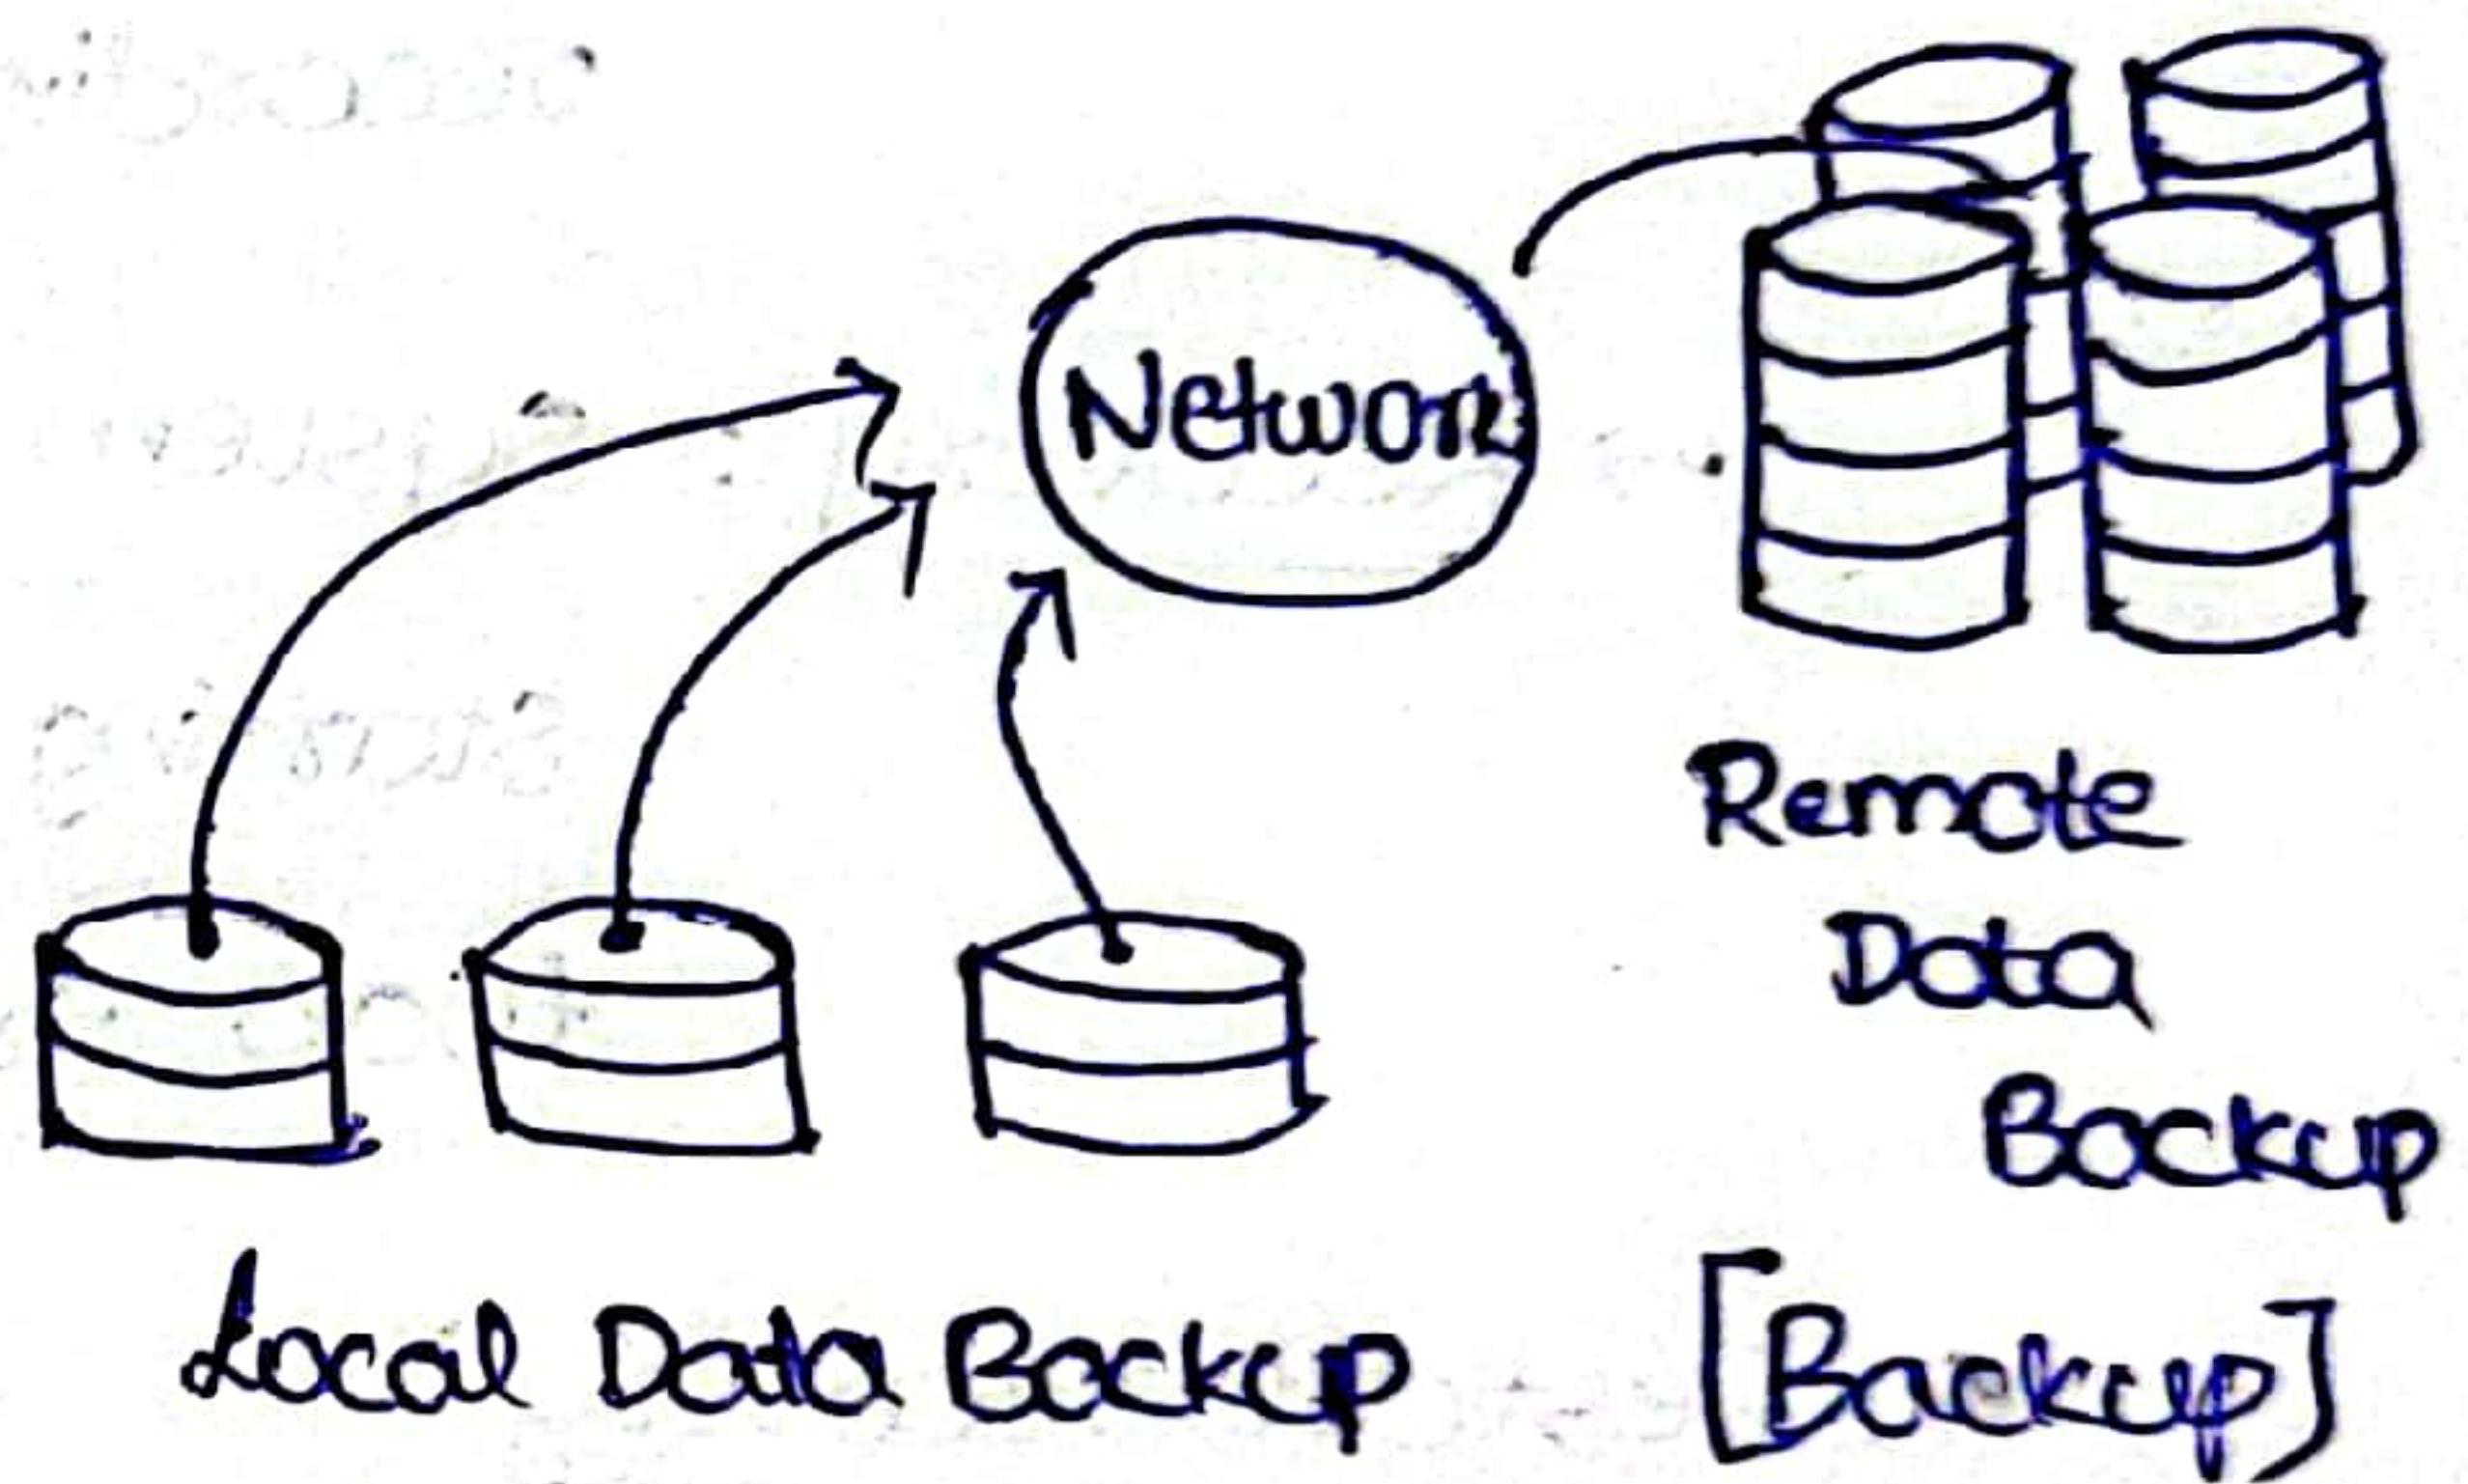
User requesting the throw-out blocks ← Buffer Manager address



Remote Backup Systems:

→ Remote backup provides a sense of security in case the primary location where the DB stored gets destroyed.

→ Remote backup can be online/offline/real-time. In case it is offline, it is maintained manually.



Transfer of Control:

- Backup site ready
- ↓
- Recovery at Backup site
- ↓
- Backup site takes over
- ↓
- Transfers control back
- ↓
- Apply updates locally
- ↓
- Old primary recovers
- ↓
- Control transferred

Time to recover:

- Backup site ready
- ↓
- Periodic recovery → continuously updating its own copy of DB from prev. state
- ↓
- checkpoint & deleting old logs
- ↓
- Hot-Spare Configuration
- ↓
- continuous redo processing
- ↓
- Detecting primary failure
- ↓
- Ready for new transactions.